

```

from matplotlib import pyplot as plt

import numpy as np
import tensorflow as tf
tf.config.run_functions_eagerly(True)
from tensorflow.keras import Sequential
from tensorflow.keras.layers import RandomFlip, RandomRotation, RandomZoom
from tensorflow.keras.layers import (
    Layer, Conv2D, Conv2DTranspose, Dense, Concatenate
)
from tensorflow.keras.layers import (
    Reshape, Flatten, Lambda
)
from tensorflow.keras.layers import (
    BatchNormalization, Dropout
)
from tensorflow.keras.layers import LeakyReLU, ReLU
from tensorflow.keras.preprocessing import image_dataset_from_directory, text_dataset_from_directory
from tensorflow.keras.utils import array_to_img
from sentence_transformers import SentenceTransformer

```

```

# GLOBAL VARIABLES (hyperparameters)
IMG_SIZE = 64
MAX_TEXT_LEN = 512
LR = 1e-4
BATCH_SIZE = 256
EPOCHS = 1000
LATENT_DIM = 100
FEATURE_MAP = 64
WEIGHT_DECAY = 1e-5

# file paths
IMAGE_PATH = "data/CUB_200_2011/images"
TEXT_PATH = "data/CUB_200_2011/text_c10"
CHK_POINT_PATH = "saved_models/gan.weights.h5"
GEN_SAVED_PATH = "saved_images"

# d_loss g_loss array
D_LOSSES, G_LOSSES = [], []

# sentences
PROMPTS = [
    'the medium sized bird has a dark grey color, a black downward curved beak, and long wings.',
    'the bird is dark grey brown with a thick curved bill and a flat shaped tail.',
    'bird has brown body feathers, white breast feathers and black beak',
    'this bird has a dark brown overall body color, with a small white patch around the base of the bill.',
    'the bird has very long and large brown wings, as well as a black body and a long black beak.',
    'it is a type of albatross with black wings, tail, back and beak, and has a white ring at the base of its beak.',
    'this bird has brown plumage and a white ring at the base of its long, curved brown beak.',
    'the entire body is dark brown, as is the bill, with a white band encircling where the bill meets the head.',
    'this bird is gray in color, with a large curved beak.',
    'a large gray bird with a long wingspan and a long black beak.',
]

```

```

def display_images(images, titles) -> None:
    """
    display images and text loaded from datasets
    """
    plt.figure(figsize=(IMG_SIZE, IMG_SIZE))
    for i in range(len(images)):
        if i > 35:
            break
        plt.subplot(6, 6, i+1)
        plt.imshow(images[i].astype('uint8'))
        plt.title(titles[i])
        plt.axis('off')
    plt.show()

def plot_loss_chart(d_losses, g_losses):
    steps = np.arange(len(d_losses))
    plt.figure(figsize=(10, 6))
    plt.plot(steps, np.array(g_losses), label='generator loss')
    plt.plot(steps, np.array(d_losses), label='discriminator loss')
    plt.xlabel('Steps')
    plt.ylabel('Loss Value')
    plt.title("Mode Collapse: Loss Trends")
    plt.legend(['Discriminator Loss', 'Generator Loss'])
    plt.show()

def plot_history_chart(history):
    plt.plot(history.history['g_loss'])
    plt.plot(history.history['d_loss'])
    plt.title('Generator and Discriminator Loss Chart')
    plt.ylabel('Losses')
    plt.xlabel('Epoch')
    plt.legend(['g_loss', 'd_loss'], loc='upper left')
    plt.show()

def normalize_images(image):
    # Convert image data type to float32 and between -1.0 to 1.0
    image = tf.cast(image, tf.float32)
    image = tf.clip_by_value((image - 127.5) / 127.5, -1.0, 1.0)

    return image

def preprocess_text_batch(text_batch):
    """
    split text in lines and remove empty lines
    choose first line, lower case and do sentence embedding
    ;return: embedding vector shape (384, )
    """

```

```

def preporcess_text_batch(text_batch):
    """
    split text in lines and remove empty lines
    choose first line, lower case and do sentence embedding
    :return: embedding vector shape (384, )
    """
    model = SentenceTransformer('sentence-transformers/all-MiniLM-L6-v2')
    def preporcess_text(text):
        lines = tf.strings.split(text, sep='\n')
        lines = tf.boolean_mask(lines, tf.strings.length(lines) > 0)
        first_line = tf.cond(tf.shape(lines)[0] > 0, lambda: lines[0], lambda: '')
        first_line = tf.strings.lower(first_line)
        # Define a Python function to decode the first line and call model.encode
        def encode_first_line(line):
            # tf.print(line.numpy().decode('utf-8'))
            return model.encode(line.numpy().decode('utf-8'))
        embedding = tf.py_function(encode_first_line, [first_line], tf.float32)
        return embedding
    # Apply preporcess_text to each text in the batch
    processed_texts = tf.map_fn(preporcess_text, text_batch, fn_output_signature=tf.float32)

    return processed_texts

def get_embedding(sentence: str):
    """
    transform a sentence to sentence embeddings vector
    :return: vector shape (384, )
    """
    model = SentenceTransformer('sentence-transformers/all-MiniLM-L6-v2')

    return model.encode([sentence])

```

```

# load image dataset from image directory
# resize them into 64 * 64 * 3
image_ds = image_dataset_from_directory(
    IMAGE_PATH,
    labels=None,
    label_mode=None,
    class_names=None,
    image_size=(IMG_SIZE, IMG_SIZE),
    batch_size=BATCH_SIZE,
    validation_split=None,
    subset=None,
    shuffle=False,
)

# load text dataset from text directory
text_ds = text_dataset_from_directory(
    TEXT_PATH,
    labels=None,
    label_mode=None,
    class_names=None,
    max_length=MAX_TEXT_LEN,
    batch_size=BATCH_SIZE,
    validation_split=None,
    subset=None,
    shuffle=False,
)

```

```

# verify that images and their captions are loaded as expected pairs
images_to_display, titles = [], []
idx = 0
for image_batch, text_batch in tf.data.Dataset.zip((image_ds, text_ds)).take(1):
    for image, text in zip(image_batch, text_batch):
        if idx > 35:
            break
        idx += 1
        # Convert the image from TensorFlow tensor to numpy array
        image_np = image.numpy()
        # Append the image and corresponding text to the list
        images_to_display.append(image_np)
        titles.append(text.numpy().decode('UTF-8'))
display_images(images_to_display, titles)

```

```

# Normalize the images between -1 to 1
image_ds = image_ds.map(normalize_images)
text_ds = text_ds.map(preporcess_text_batch)

# text_ds = text_ds.map(lambda x: tf.expand_dims(x, axis=-1))

# combine image and text dataset so that we can shuffle it asynchronously
combined_ds = tf.data.Dataset.zip((image_ds, text_ds))
combined_ds = combined_ds.shuffle(buffer_size=combined_ds.cardinality()) \
    .prefetch(buffer_size=tf.data.AUTOTUNE)

```

```

for image, text in combined_ds.take(1):
    print(tf.shape(image))
    print('image', image.numpy())
    print('text', text.numpy())
    break

```

```

class AugmentLayer(Layer):
    """
    Data augmentation layer for image input
    """

    def __init__(self, name: str = 'AugmentationLayer', **kwargs):
        super(AugmentLayer, self).__init__(name=name, **kwargs)

        # create a sequential block
        self.augment_layer = Sequential([
            RandomFlip("horizontal"),
            RandomRotation(0.1),
            RandomZoom(0.2)
        ], name=name)

    def call(self, inputs, training=False, mask=None):
        x = self.augment_layer(inputs)

        return x

```

```

class ReshapeLayer(Layer):
    """
    Reshape layer for embedding vector
    """

    def call(self, inputs, training=False, mask=None):
        # Reshape the tensor
        text = tf.reshape(inputs, (-1, 1, 1, 128))

        # Tile the tensor
        text = tf.tile(text, (1, 4, 4, 1))

        return text

```

```

class ConditioningAugmentationLayer(Layer):
    """
    This layer basically tokenizes text input and does word embeddings
    it returns conditional variated embedded vector (128-d vector)
    """

    def __init__(self, embed_dim: int = 256,
                 name: str = 'Conditionallayer',
                 **kwargs):
        super(ConditioningAugmentationLayer, self).__init__(name=name, **kwargs)
        self.embed_dim = embed_dim

        self.dense_layer = Dense(self.embed_dim)
        self.act_layer = LeakyReLU(negative_slope=0.2)
        self.lambda_layer = Lambda(self.conditioning_augmentation)

    def call(self, inputs, training=False, mask=None):

        x = self.dense_layer(inputs) # return shape(embed_dim, )
        x = self.act_layer(x) # return shape(embed_dim, )
        x = self.lambda_layer(x)

        return x

    def compute_output_shape(self, input_shape):
        # Calculate the output shape based on input_shape and layer configurations
        # (e.g., embedding size, sequence length after processing)
        output_shape = tf.TensorShape([input_shape[0], self.embed_dim//2 ])
        return output_shape

    def conditioning_augmentation(self, input):
        """

```

```

        This function takes the output of the text embedding (denoted as input)
        and splits it into two parts: mean and log_sigma.
        'mean' represents the mean of the Gaussian distribution,
        'log_sigma' represents the logarithm of the standard deviation.
        :returns: conditioning augmentation vector
        """

        # conditioning augmentation process
        mean = input[:, :self.embed_dim//2]
        log_sigma = input[:, self.embed_dim//2:]
        stddev = tf.math.exp(log_sigma)
        epsilon = tf.random.normal(shape=tf.cast(tf.shape(mean), tf.int32), dtype=mean.dtype)
        # epsilon = tf.random.normal(shape=(tf.shape(mean)[0], tf.shape(mean)[1], self.embed_dim), dtype=mean.dtype)
        c = mean + stddev * epsilon # return shape(embed_dim / 2, )

        return c

```

```

class EncoderLayer(Layer):
    """
    create basic encoder layer blocks for constructing
    nn architecture of discriminator model
    """

    def __init__(self, filters: int, kernel: int = 3, bn: bool = True, name=None):
        super(EncoderLayer, self).__init__(name=name)
        self.filters = filters
        self.kernel = kernel
        self.bn = bn

        # using convolution layer to downsample the image
        self.conv_layer = Conv2D(filters, kernel_size=4,
                                   strides=2,
                                   padding='same',
                                   use_bias=False,
                                   kernel_regularizer='l2',
                                   kernel_initializer='he_uniform')

        # using LeakyRelu as activation function
        self.act_layer = LeakyReLU(negative_slope=0.2)
        if self.bn:
            # using batch normalization to regularize the mode
            self.bn_layer = BatchNormalization(momentum=0.8)

    def call(self, inputs, training=False, mask=None):
        x = self.conv_layer(inputs)
        x = self.act_layer(x)
        if self.bn:
            x = self.bn_layer(x, training=training)

        return x

```

```

class DecoderLayer(Layer):
    """
    create basic decoder layer blocks for constructing
    nn architecture of generator model
    """

    def __init__(self, filters: int, kernel: int = 3, bn: bool = True, name=None):
        super(DecoderLayer, self).__init__(name=name)
        self.filters = filters
        self.kernel = kernel
        self.bn = bn

        # using convolution layer to upsample the image
        self.deconv_layer = Conv2DTranspose(filters,
                                              kernel_size=4,
                                              strides=2,
                                              padding='same',
                                              use_bias=False,
                                              kernel_regularizer=tf.keras.regularizers.l2(0.01),
                                              kernel_initializer='he_uniform')

        if self.bn:
            # using batch normalization to regularize the mode
            self.bn_layer = BatchNormalization(momentum=0.8)
        # using LeakyRelu as activation function
        self.act_layer = ReLU()
        # weaken the discriminator
        self.dropout_layer = Dropout(0.3)

    def call(self, inputs, training=False, mask=None):
        x = self.deconv_layer(inputs)
        if self.bn:
            x = self.bn_layer(x, training=training)
        x = self.act_layer(x)
        x = self.dropout_layer(x)

        return x

```

```

class Discriminator(tf.keras.Model):
    """
    the discriminator is fed with two inputs, the image and text pair from Generator
    and returns True or False
    """
    def __init__(self, input_shape,
                  latent_dim: int = 100,
                  feature_maps: int = 64,
                  name='DiscriminatorModel',
                  **kwargs):
        """
        initialize the DISCRIMINATOR model of gan
        :param input_shape: (h, w, channels)
        :param feature_maps:
        :param kwargs:
        """
        super(Discriminator, self).__init__(name=name, **kwargs)
        self.input_shape = input_shape
        self.latent_dim = latent_dim
        self.feature_maps = feature_maps

        # Data augmentation
        self.augment_layer = AugmentLayer(name='AugmentLayer')

        # encoder (downsampling) layers to compress information
        self.conv_layer1 = EncoderLayer(filters=self.feature_maps, name=f'EncConv{self.feature_maps}')
        self.conv_layer2 = EncoderLayer(filters=self.feature_maps * 2, name=f'EncConv{self.feature_maps * 2}')
        self.conv_layer3 = EncoderLayer(filters=self.feature_maps * 4, name=f'EncConv{self.feature_maps * 4}')
        self.conv_layer4 = EncoderLayer(filters=self.feature_maps * 8, name=f'EncConv{self.feature_maps * 8}')

        # text embedding processing layer: text embedding -> (128, )
        self.ca_layer = ConditioningAugmentationLayer()
        self.ca_reshape_layer = ReshapeLayer()

        # concat image input with embedding input
        self.concat_layer = Concatenate(axis=-1, name='ConcatLayer')

        # change the last rank from 640 to 512 after concatenate
        self.conv_kernel1_layer = Conv2D(self.feature_maps * 8,
                                         kernel_size=1,
                                         padding='same',
                                         use_bias=False,
                                         kernel_regularizer=tf.keras.regularizers.l2(0.01),
                                         kernel_initializer='he_uniform')
        self.conv_kernel1_bn_layer = BatchNormalization()
        self.conv_kernel1_act_layer = LeakyReLU(negative_slope=0.2)

        # flatten data into 1D
        self.flat_layer = Flatten(name='FlattenLayer')
        # add regularization layer
        self.drop_layer = Dropout(0.25, name='DropoutLayer')

```

```

# add regularization layer
self.drop_layer = Dropout(0.25, name='DropoutLayer')

self.out_layer = Dense(1, activation='sigmoid', name='ClsOutputLayer')

def call(self, inputs, training=False, mask=None):
    image, text = inputs

    # data augmentation -> (64, 64, 3)
    image = self.augment_layer(image)

    # forward through encoder layers -> (4, 4, 512)
    x = self.conv_layer1(image, training=training)
    x = self.conv_layer2(x, training=training)
    x = self.conv_layer3(x, training=training)
    x = self.conv_layer4(x, training=training)

    text = self.ca_layer(text)
    text = self.ca_reshape_layer(text)

    concat = self.concat_layer([x, text])

    # convert concatenated shape into (4, 4, 512)
    xx = self.conv_kernel1_layer(concat)
    xx = self.conv_kernel1_bn_layer(xx)
    xx = self.conv_kernel1_act_layer(xx)

    # add reshape and regularization layers
    xx = self.flat_layer(xx)
    xx = self.drop_layer(xx, training=training)

    return self.out_layer(xx)

def build_graph(self):
    txt = tf.keras.Input(shape=(384, ), name='EmbeddingInputLayer')
    img = tf.keras.Input(shape=self.input_shape, name='ImageInputlayer')

    return tf.keras.Model(inputs=[img, txt], outputs=self.call([img, txt]))

```



```

class Generator(tf.keras.Model):
    """
    generator takes two inputs: text and latent vector
    and returns 64 * 64 * 3 image and text pair
    """
    def __init__(self, latent_dim: int = 100,
                  feature_maps: int = 64,
                  name='GeneratorModel',
                  **kwargs):
        """
        initialize the GENERATOR model of gan
        :param latent_dim: int
        :param feature_maps:
        :param name: str
        :param kwargs:
        """
        super(Generator, self).__init__(name=name, **kwargs)
        self.latent_dim = latent_dim
        self.feature_maps = feature_maps

        # text embedding processing layer: text embedding (384, )
        self.ca_layer = ConditioningAugmentationLayer()
        self.concat_layer = Concatenate(axis=1, name='ConcatLayer')

        # fully connected layer and reshape layer to process input latent vector
        self.dense_layer = Dense(256 * 8 * 8, use_bias=False, activation='relu', name='FCLayer')
        self.reshape_layer = Reshape((4, 4, -1), name='ReshapeLayer')

        # decode layers to de-conv (upsampling) layers and generate image
        self.conv_layer1 = DecoderLayer(filters=self.feature_maps * 8, name=f'DecConv{self.feature_maps * 8}')
        self.conv_layer2 = DecoderLayer(filters=self.feature_maps * 4, name=f'DecConv{self.feature_maps * 4}')
        self.conv_layer3 = DecoderLayer(filters=self.feature_maps * 2, name=f'DecConv{self.feature_maps * 2}')
        self.conv_layer4 = DecoderLayer(filters=self.feature_maps, name=f'DecConv{self.feature_maps}')

```

```

        # generate image
        self.out_img_layer = Conv2D(filters=3,
                                     kernel_size=3,
                                     padding='same',
                                     activation='tanh',
                                     kernel_initializer='he_uniform',
                                     use_bias=False,
                                     name='ImageOutputLayer')

    def call(self, inputs, training=False, mask=None):
        text, noise = inputs

        text_embedding = self.ca_layer(text)

        x = self.concat_layer([text_embedding, noise])

        x = self.dense_layer(x)
        x = self.reshape_layer(x)

        # forward through encoder layers
        x = self.conv_layer1(x, training=training)
        x = self.conv_layer2(x, training=training)
        x = self.conv_layer3(x, training=training)
        x = self.conv_layer4(x, training=training)
        x1 = self.out_img_layer(x)

        return x1, text

    def build_graph(self):
        text = tf.keras.Input(shape=(384, ), name='EmbeddingInputLayer')
        noise = tf.keras.Input(shape=(self.latent_dim, ), name='NoiseInputLayer')

        return tf.keras.Model(inputs=[text, noise], outputs=self.call([text, noise]))

```

```

class GanModel(tf.keras.Model):
    """
    create GAN model by implementing both generator and discriminator in the model
    1. generate random noise vector
    2. create fake images using generator and noise vector
    3. concatenate both fake images (0) and real images (1) with labels
    4. train discriminator model with above dataset
    5. generate another noise vector and fake images (1) for generator model training
    """

    def __init__(self, disc_model, gen_model,
                 latent_dim: int = 128,
                 name='GanModel', **kwargs):
        """
        :param disc_model: discriminator model
        :param gen_model: generator model
        :param latent_dim: latent vector dimensions
        """
        super(GanModel, self).__init__(name=name, **kwargs)
        self.disc_model = disc_model
        self.gen_model = gen_model
        self.latent_dim = latent_dim

    def compile(self, d_optimizer, g_optimizer, loss_fn_d, loss_fn_g):
        """
        setup compile configurations
        :param d_optimizer: optimizer for discriminator
        :param g_optimizer: optimizer for generator
        :param loss_fn_g: loss function for generator model
        :param loss_fn_d: loss function for discriminator model
        :return:
        """
        super(GanModel, self).compile()
        # setup optimizer for discriminator and generator respectively
        self.d_optimizer = d_optimizer
        self.g_optimizer = g_optimizer
        # setup the loss function for the model, binary_crossentropy should be good
        self.loss_fn_d = loss_fn_d
        self.loss_fn_g = loss_fn_g
        # setup metrics for both discriminator model and generator model respectively
        # using binary cross entropy as metrics because this is a binary classification task
        self.d_loss_metric = tf.keras.metrics.BinaryCrossentropy(name="d_loss")
        self.g_loss_metric = tf.keras.metrics.Mean(name="g_loss")

    @property
    def metrics(self):
        return [self.d_loss_metric, self.g_loss_metric]

    @tf.function
    def train_step(self, data):
        """
        take the real images and generate fake images, each step processes a batch size of image

```

```

@tf.function
def train_step(self, data):
    """
    take the real images and generate fake images, each step processes a batch size of image
    :param data: real images and text from dataset
    :return:
    """
    real_images, real_text = data

    # batch_size = BATCH_SIZE
    batch_size = tf.shape(real_images)[0]

    # generate random noise vector to generate fake images for discriminator training
    latent_vectors_disc = tf.random.normal(
        shape=(batch_size, self.latent_dim)
    )

    # generate fake images from real text and noise vector
    fake_images_disc, fake_text = self.gen_model([real_text, latent_vectors_disc])

    # Combine them with real images and labels
    combined_images_disc = tf.concat([fake_images_disc, real_images], axis=0)
    combined_text_disc = tf.concat([fake_text, real_text], axis=0)

    labels_disc = tf.concat(
        [tf.ones((batch_size, 1)) * 0.95, tf.zeros((batch_size, 1))],
        axis=0
    )

    # Add random noise to the all labels
    labels_disc += tf.random.uniform(tf.shape(labels_disc)) * 0.05

    # Train the discriminator
    with tf.GradientTape() as tape:
        # get predictions from discriminator model
        d_predictions = self.disc_model([combined_images_disc, combined_text_disc])
        # calculate the loss from discriminator model
        d_loss = self.loss_fn_d(labels_disc, d_predictions)

    # get gradient of the discriminator model and update weights
    grads_disc = tape.gradient(d_loss, self.disc_model.trainable_weights)
    self.d_optimizer.apply_gradients(
        zip(grads_disc, self.disc_model.trainable_weights)
    )

    # generate random noise vector to generate fake images for generator training
    latent_vectors_gen = tf.random.normal(
        shape=(batch_size, self.latent_dim)
    )

    # Assemble labels that say "all real images"
    labels_gen = tf.zeros((batch_size, 1))

```

```

# generate random noise vector to generate fake images for generator training
latent_vectors_gen = tf.random.normal(
    shape=(batch_size, self.latent_dim)
)

# Assemble labels that say "all real images"
labels_gen = tf.zeros((batch_size, 1))

# Train the generator (note that we should *not* update the weights
# of the discriminator)!
with tf.GradientTape() as tape:
    # generate another fake images for discriminator
    gen_fake_images, gen_text = self.gen_model([fake_text, latent_vectors_gen])
    # generate fake images and get predictions from discriminator model
    predictions = self.disc_model([gen_fake_images, gen_text])
    # calculate the loss for generator model
    g_loss = self.loss_fn_g(labels_gen, predictions)

# Print discriminator and generator losses
tf.print("d_loss:", d_loss, "g_loss:", g_loss)

# get the gradients for generator model and update generator weights
grads_gen = tape.gradient(g_loss, self.gen_model.trainable_weights)
self.g_optimizer.apply_gradients(
    zip(grads_gen, self.gen_model.trainable_weights)
)

# Update metrics
self.d_loss_metric.update_state(labels_disc, d_predictions)
self.g_loss_metric.update_state(g_loss)

G_LOSSES.append(g_loss)
D_LOSSES.append(d_loss)

return {
    "d_loss": self.d_loss_metric.result(),
    "g_loss": self.g_loss_metric.result(),
}

```

```

class GANMonitor(tf.keras.callbacks.Callback):
    def __init__(self, num_img=3, latent_dim=100):
        self.num_img = num_img if num_img <= 5 else 5
        self.latent_dim = latent_dim
        self.prompt_embedding = tf.reshape(tf.constant(list(map(get_embedding, PROMPTS[:self.num_img]))), (self.num_img, -1))

    def on_epoch_end(self, epoch, logs=None):
        random_latent_vectors = tf.random.normal(
            shape=(self.num_img, self.latent_dim)
        )
        generated_images, _ = self.model.gen_model([self.prompt_embedding,
                                                    random_latent_vectors])
        generated_images = (generated_images + 1) * 127.5
        for i in range(self.num_img):
            img = array_to_img(generated_images[i])
            img.save("%sgenerated_img_%03d_%d.png" % (GEN_SAVED_PATH, epoch, i))

```

```

# Instantiate the ConditioningAugmentationLayer
conditioning_layer = ConditioningAugmentationLayer()
# Provide sample input data (text embeddings)
# Assuming inputs is a numpy array of shape (batch_size, sequence_length)
# Call the layer with the input data
for text in text_ds.take(1):
    output = conditioning_layer(text)
    # Inspect the output
    print(output.shape)
    break

```

```

# Create an instance of the Discriminator model
generator = Generator(latent_dim=100, feature_maps=64)
# view the summary of the generator model
generator.build_graph().summary()

```

```

discriminator = Discriminator(input_shape=(64, 64, 3), latent_dim=128, feature_maps=64)

# view the summary of the generator model
discriminator.build_graph().summary()

```

```

# callback functions
# early stopping callback configure
cb_early_stop = tf.keras.callbacks.EarlyStopping(
    monitor="g_loss",
    patience=50,
    verbose=1,
    mode="min",
    restore_best_weights=True,
)

cb_check_point = tf.keras.callbacks.ModelCheckpoint(
    CHK_POINT_PATH,
    monitor='g_loss',
    verbose=1,
    save_best_only=True,
    save_weights_only=True,
    mode='min',
    save_freq='epoch',
)

```

```

# Initialize GanModel by providing the discriminator model, generator model and latent dim
# discriminator.trainable = False
gan = GanModel(disc_model=discriminator,
               gen_model=generator,
               latent_dim=LATENT_DIM)

# configure the gan optimizers and loss function
# discriminator take lower learning rate and generator takes higher learning rate
# so that generator is not overwhelmed by discriminator
gan.compile(
    d_optimizer=tf.keras.optimizers.Adam(learning_rate=LR/2, weight_decay=WEIGHT_DECAY),
    g_optimizer=tf.keras.optimizers.Adam(learning_rate=LR, weight_decay=WEIGHT_DECAY),
    loss_fn_d=tf.keras.losses.BinaryCrossentropy(from_logits=False),
    loss_fn_g=tf.keras.losses.MeanSquaredError(),
)

```

```
# train the model
history = gan.fit(
    combined_ds,
    epochs=EPOCHS,
    # validation_split=0.1,
    verbose=2,
    callbacks=[cb_early_stop,
               cb_check_point,
               GANMonitor(num_img=10, latent_dim=LATENT_DIM)
              ],
)
```

```
# display the loss chart of gan model
plot_loss_chart(d_losses=D_LOSSES, g_losses=G_LOSSES)
```

```
plot_history_chart(history)
```

```
prompt = get_embedding('A bird is brown with a lighter brown crest')
latent_vector = tf.random.normal(shape=(1, LATENT_DIM))
image, _ = gan.gen_model([prompt, latent_vector])
```